

▼ COGS 182 Project 2 | Schedule Optimization

▼ Checkpoint 3: Implementing the Algorithms

```
import numpy as np
import matplotlib.pyplot as plt
import sys
from mpl_toolkits import mplot3d
import seaborn as sns
from tqdm.auto import tqdm
import pickle
```

▼ List of Possible Topics & Actions (Hour Allocations to Each Topic) & Resultant States

```
possible_topics=['topology',
                 'Japanese',
                 'webtooning',
                 'physics',
                 'video-editing',
                 'graphic design',
                 'animation',
                 'psych/linguistics',
                 'R/statistical packages',
                 'front-end coding']
```

```
#0-12 = allot hrs, 7 = commit to schedule
possible_actions = [0,1,3,5,7]
```

```
r_times = []
for remaining_time in np.arange(25):
    for action in [0,1,3,5]:
        for org_hrs in [0,1,2,3,4]:
            for lab_hrs in [0, 1,2]:
                r_time = remaining_time - action - org_hrs - lab_hrs
                r_times.append(r_time)
```

```
possible_remaining_times = sorted(np.unique(r_times))
```

```
terminal_states = []
for r_time in possible_remaining_times:
    terminal_state = [r_time, 0]
    terminal_states.append(terminal_state)
```

```
possible_states = [0]
for r_time in possible_remaining_times:
    for len_schedule in np.arange(len(possible_topics) +1):
        for topic in possible_topics:
            state = [r_time, len_schedule, topic]
            if state[0] > 0:
                possible_states.append(state)
```

```
# len(terminal_states),len(possible_states)
```

```
pos_states = [0]
for state_indx, state in enumerate(possible_states[1:]):
    r_time, len_schd, topic = state
    no_topic_state = [r_time, len_schd]
    if no_topic_state not in pos_states:
        pos_states.append(no_topic_state)
```

▼ Factors (i.e. schedule fulfillment, coherence) to take into account when calculating the rewards

```
#mutually reinforcing subjects (not really accurate to real life, but it suffices)
related_topics = {'topology': ['physics', 'R/statistical packages', 'graphic design'],
                  'Japanese': ['linguistics', 'animation', 'webtooning'],
                  'webtooning': ['animation', 'Japanese', 'graphic design'],
                  'physics': ['topology','chemistry', 'R/statistical packages'],
                  'video-editing': ['animation', 'content creation', 'webtooning'],
                  'graphic design': ['webtooning', 'front-end coding'],
                  'animation': ['webtooning', 'video-editing', 'graphic design', 'content creation'],
                  'psych/linguistics': ['Japanese', 'R/statistical packages'],
                  'R/statistical packages': ['educational psychology', 'physics'],
                  'front-end coding': ['graphic design', 'animation', 'content creation']}
```

```
# personal ratings of [frustration/learning curve, intrigue, applicability] by topic
```

```

fulfillment_factors = {'topology': [1, 1, 0.6],
                      'Japanese': [0.7, 1, 1],
                      'webtooning': [0.5, 0.5, 0.4],
                      'physics': [1, 0.9, .6],
                      'video-editing': [0.8, 0.4, 1],
                      'graphic design': [0.5, 0.3, 0.6],
                      'animation': [1, 0.9, 0.7],
                      'psych/linguistics': [0.6, 0.7, 0.7],
                      'R/statistical packages': [0.5, 0.4, 1],
                      'front-end coding': [0.6, 0.7, 1]}

fulfillment_scores = {}
for topic in fulfillment_factors:
    frustration, intrigue, applicability = fulfillment_factors[topic]

    #fulfillment factors scale positive rewards
    fulfillment_score = (-0.25*frustration + 0.5*intrigue + 0.25*applicability)
    fulfillment_scores[topic] = fulfillment_score

```

▼ Step function, take in remaining time (in a day), the schedule, the topic of consideration, and the action

- NOTE: for feasibility sake, the observed state is actually only the LENGTH of the current schedule, but the entire schedule is passed for the environment to do calculations

```

# state keeps track of remaining free time in a day, the length of the schedule, and the current topic
# length of schedule instead of actual schedule, since that's too many different states...

def step(remaining_time, schedule, topic, action):

    state = [remaining_time, len(schedule)]
    state_indx = pos_states.index(state)

    if action != 7:
        if action != 0:
            remaining_time -= action
            schedule[topic] = action #then allot hours according to a policy?

        topic = np.random.choice(possible_topics)
        new_state = [remaining_time, len(schedule)]
        reward = 0

        if remaining_time < 6:
            # print("\033[1;31m  KAROSHI  \033[0m YOU DIED OF OVERWORK")
            # sys.stdout.flush()
            reward = -1
            new_state = 0

        return new_state, reward, schedule, topic

# ENVIRONMENT GOES when commit to schedule
elif action == 7:
    new_state, reward, schedule, remaining_time = env(remaining_time, schedule)

    return new_state, reward, schedule, remaining_time

```

▼ When the environment goes,

it takes the remaining time and schedule, adds on random hours from existing commitments, calculates if the remaining time in a day allows for sleep, calculates the schedule fulfillment and modulates the fulfillment with bonuses for schedule coherence and having enough to sleep. The terminal state is returned [remaining_time, 0].

```

def env(remaining_time, schedule):

    existing_commitments = {'Orgs': np.random.choice(np.arange(1,5)), 'LabStuff': np.random.choice(np.arange(1,3))}
    # factor in existing commitments
    remaining_time -= (existing_commitments['Orgs'] + existing_commitments['LabStuff'])

    # KAROSHI AGAIN if forget about current commitments
    if remaining_time < 6:
        reward = -1

    #penalty for undercommitment (listlessness)
    elif remaining_time > 18:
        reward = -0.5

    else:
        total_fulfillment = 0

        # get average of fulfillments of all commitments in schedule
        for topic in schedule:
            if topic in ["Orgs", "LabStuff"]:
                continue

```

```

        else:
            total_fulfillment += fulfillment_scores[topic]
        total_fulfillment = np.mean(total_fulfillment)

    # factor in related subjects into fulfillment score (mutually reinforcing)
    # if too many subjects at once, coherence turns into distraction factor, negative rate
    coherence = 1
    if len(schedule) >= 3:
        for topic in schedule:
            for possible_topic in possible_topics:
                if topic == possible_topic:
                    for related_topic in related_topics[topic]:
                        if related_topic in schedule:
                            coherence += 0.05
    elif len(schedule) > 5: #too many topics
        coherence = -0.5

    # sleep and free time bonus (a rate)
    if remaining_time >= 8 and remaining_time <= 10:
        bonus = 1.1          # 10% bonus fulfillment
    else:
        bonus = 1            # no bonus

    # reward to return as a function of bonus, topic-based fulfillment and coherence of schedule
    reward = (total_fulfillment)* ((bonus + coherence)/ 2)

    schedule['Orgs'] = existing_commitments['Orgs']
    schedule['LabStuff'] = existing_commitments['LabStuff']

    new_state = 0

    return new_state, reward, schedule, remaining_time

```

▼ Algorithm 1: On-Policy First-Visit Monte Carlo Control

```

def init_q_values(init):
    if init == "zeros":
        q_values = np.zeros((len(pos_states), len(possible_actions)))
    elif init == 'arb':
        q_values = np.ones((len(pos_states), len(possible_actions)))
        for indx, value in enumerate(q_values):
            q_values[indx] = np.random.rand()
    else:
        q_values = np.ones((len(pos_states), len(possible_actions))) * init
    return q_values

def MonteCarlo(init, num_runs):
    # num of times that action a has been selected from state s
    N = np.zeros((len(pos_states), len(possible_actions)))
    N_0 = 1

    # hyperparameters
    gamma = 0.5    # discount factor - balance between being too myopic (0) and too farsighted (1)
    epsilon = N_0/(N_0 + 0)

    # initialize policy, q_values, returns

    q_values = init_q_values(init)
    returns = np.zeros((len(pos_states), len(possible_actions)))
    policy = np.ones((len(pos_states), len(possible_actions))) * (epsilon/len(possible_actions))

    all_episodes = []
    all_episode_rewards = []

    for run in tqdm((range(num_runs)), position=0):

        remaining_time = 24
        schedule = {}
        topic = np.random.choice(possible_topics)

        state      = [remaining_time, len(schedule)]
        state_indx = pos_states.index(state)

        action = np.random.choice(possible_actions, p= policy[state_indx])

        episode = []
        episode_rewards = []

        while True:
            if action != 7:

```

```

        new_state, reward, schedule, topic = step(remaining_time, schedule, topic, action)
    elif action == 7:
        new_state, reward, schedule, remaining_time = step(remaining_time, schedule, topic, action)

    ep = [state, action, reward, schedule]
    episode.append(ep)
    episode_rewards.append(reward)

    if new_state == 0:
        break

    remaining_time, len_schedule = new_state
    state = [remaining_time, len_schedule]
    state_indx = pos_states.index(state)

    action = np.random.choice(possible_actions, p= policy[state_indx])
    action_indx = possible_actions.index(action)

all_episodes.append(episode)
all_episode_rewards.append(episode_rewards)

G = 0
#loop through episode in reverse
for indx_rev, state_action in enumerate(episode[::-1]):
    state, action, reward, schedule = state_action
    state_indx = pos_states.index(state)
    action_indx = possible_actions.index(action)

    indx_state_action = len(episode_rewards) - indx_rev -1

    G += gamma*G + episode_rewards[indx_state_action]
    returns[state_indx][action_indx] = G

    if state_action not in episode[:indx_state_action]:
        N[state_indx][action_indx] += 1
        alpha = 1/(N[state_indx][action_indx]) # use time-varying alpha
        q_values[state_indx][action_indx] += alpha * (G - q_values[state_indx][action_indx])
        epsilon = N_0/(N_0 + np.min(N[state_indx][action_indx])) # epsilon-greedy exploration strategy

        #update policy
        optimal_action = np.argmax(q_values[state_indx])
        for action_indx in range(len(possible_actions)):
            if action_indx == optimal_action:
                policy[state_indx][action_indx] = 1 - epsilon + (epsilon/ len(possible_actions))
            elif action_indx != optimal_action:
                policy[state_indx][action_indx] = (epsilon/ len(possible_actions))

return policy, q_values, all_episodes, all_episode_rewards

```

```

def quick_check(q_values, policy):
    #quick check of Q-values vs. policy

    print("<<QUICK CHECKS>>")
    oops = [100, 150, 200, 250]

    for oop in oops:
        print("-----")
        print("state:", pos_states[oop], "\nQs:", q_values[oop], "\nPi:", policy[oop])
        print("Q says", np.argmax(q_values[oop]), "| Pi says", np.argmax(policy[oop]))

    agreements = 0
    for state_indx, state in enumerate(pos_states):
        if np.argmax(q_values[state_indx]) == np.argmax(policy[state_indx]):
            agreements += 1
    print("Policy aligns with Q-values:", agreements / len(pos_states) * 100, "%")

    sys.stdout.flush()

```

```

def get_opt_actions(policy):

    if algo == "MC":
        optimal_actions = []
        for state_indx, state in enumerate(pos_states):
            action_indx = np.argmax(policy[state_indx])
            if action_indx == 0:
                action = 0
                optimal_actions.append(action)
            elif action_indx == 1:
                action = 1
                optimal_actions.append(action)
            elif action_indx == 2:
                action = 3
                optimal_actions.append(action)
            elif action_indx == 3:
                action = 5

```



```
return optimal_schedules, part_optimal_schedules, optimal_schd_rewards, part_opt_schd_rewards
```

▼ Run Monte Carlo using different Initializations

```
%%time

num_runs = 100000
#try with different initialization of q_values
inits = [-0.1, 0.1,'zeros', 'arb']

policies = []
MC_init_qs = []
MC_init_eps = []
MC_init_ep_rewards = []
MC_init_opt_actions = []

for init in inits:
    print("\033[1;31mINIT Q-VALUES AS {}: \033[0m\n".format(init))

    policy, q_values, all_episodes, all_episode_rewards = MonteCarlo(init, num_runs)
    optimal_actions = get_opt_actions(policy)

    policies.append(policies)
    MC_init_qs.append(q_values)
    MC_init_eps.append(all_episodes)
    MC_init_ep_rewards.append(all_episode_rewards)
    MC_init_opt_actions.append(optimal_actions)
    print()

#    quick_check(q_values,policy)
#    print()
#    for action in possible_actions:
#        if action != 7:
#            print("optimal to allot",action, "hrs:", optimal_actions.count(action), "x")
#        else:
#            print("Optimal to commit:      ", optimal_actions.count(action), "x")
#    print()

#calculate V_star & Plots
V_star = calc_V_star(q_values)
plot_values_actions_states("MC", q_values)
plot_heatmap(V_star)
best_r_time, best_len = pos_states[np.argmax(V_star)]
print("PRIOR TO ORGS/LABSTUFF:")
print("Optimal Remaining Time: {} HOURS \nOptimal Schedule Length: {} COMMITMENTS".format(best_r_time, best_len))
print()

# returns schedules that are EITHER the optimal length or optimal, get unique entries
optimal_schedules, part_optimal_schds, opt_schd_rewards, part_opt_rewards = get_opt_schedules(V_star, all_episodes, optimal_actions)
optimal_schedules = [dict(entry) for entry in set(frozenset(schd.items()) for schd in optimal_schedules)]
part_optimal_schds = [dict(entry) for entry in set(frozenset(schd.items()) for schd in part_optimal_schds)]

print("WITH ORGS/LABS IN FINAL CONSIDERATION")
print("Optimal Schedules:")
for indx, schd in enumerate(optimal_schedules):

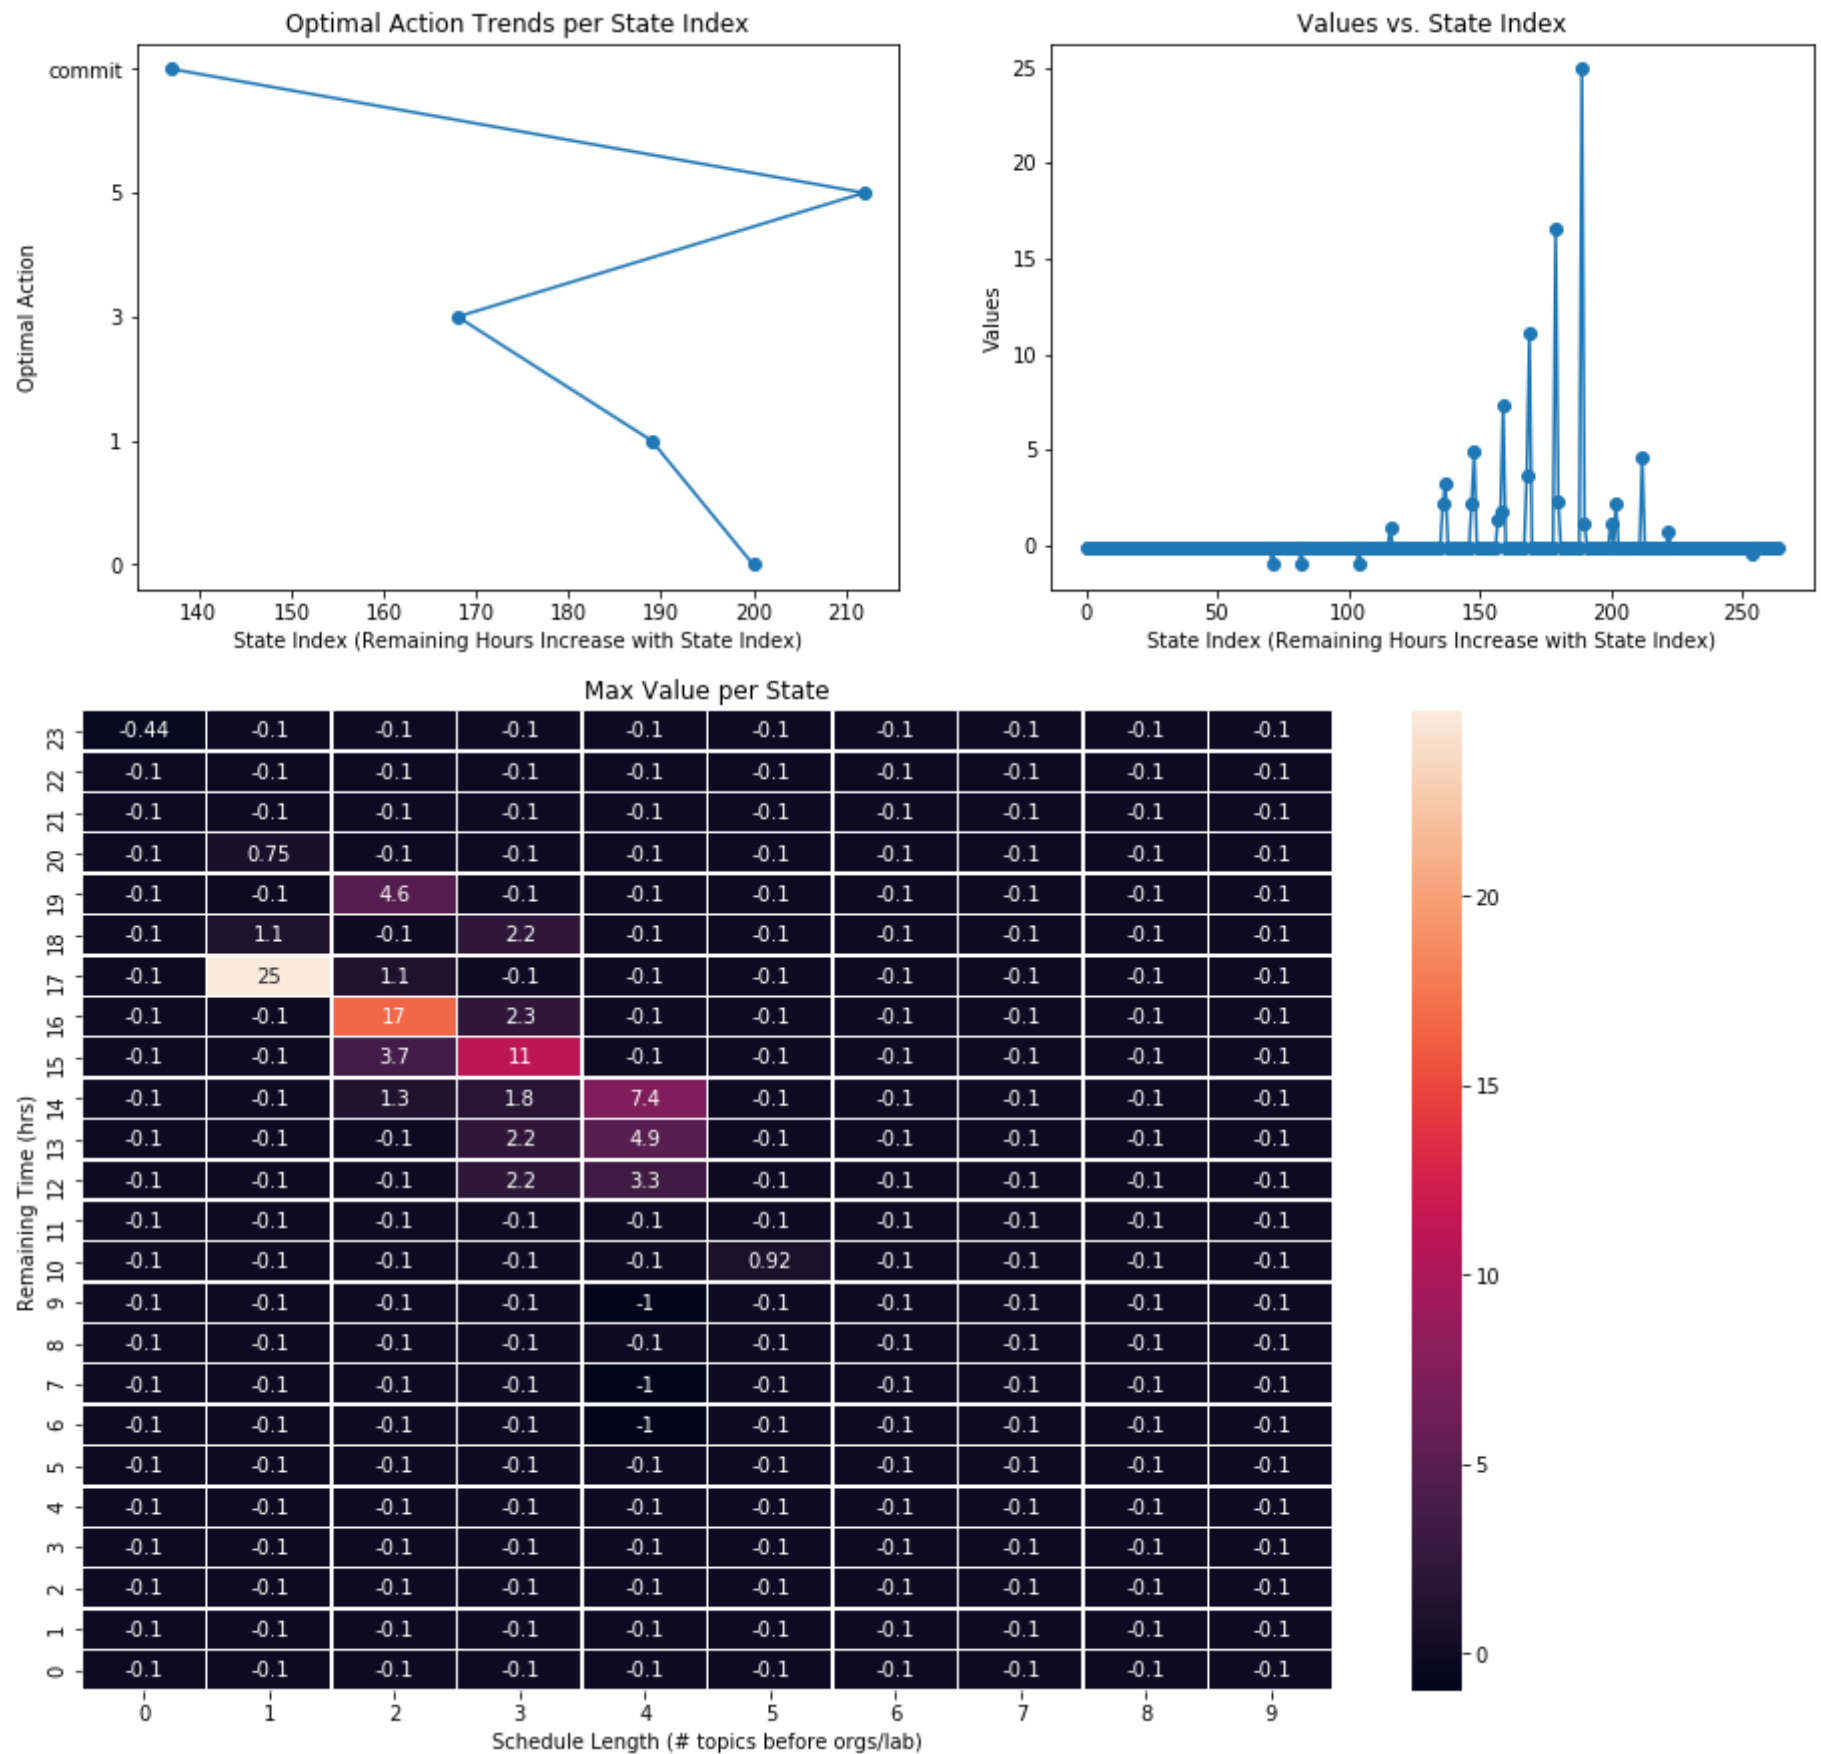
    print(schd, "| Fulfillment:", opt_schd_rewards[indx])

if optimal_schedules == []:
    print("Partially Optimal Schedules:")
    for indx, schd in enumerate(part_optimal_schds):
        print(schd, "| Fulfillment:", part_opt_rewards[indx])

print("-----\n")
sys.stdout.flush()
```

INIT Q-VALUES AS -0.1:

HBox(children=(IntProgress(value=0, max=100000), HTML(value='')))




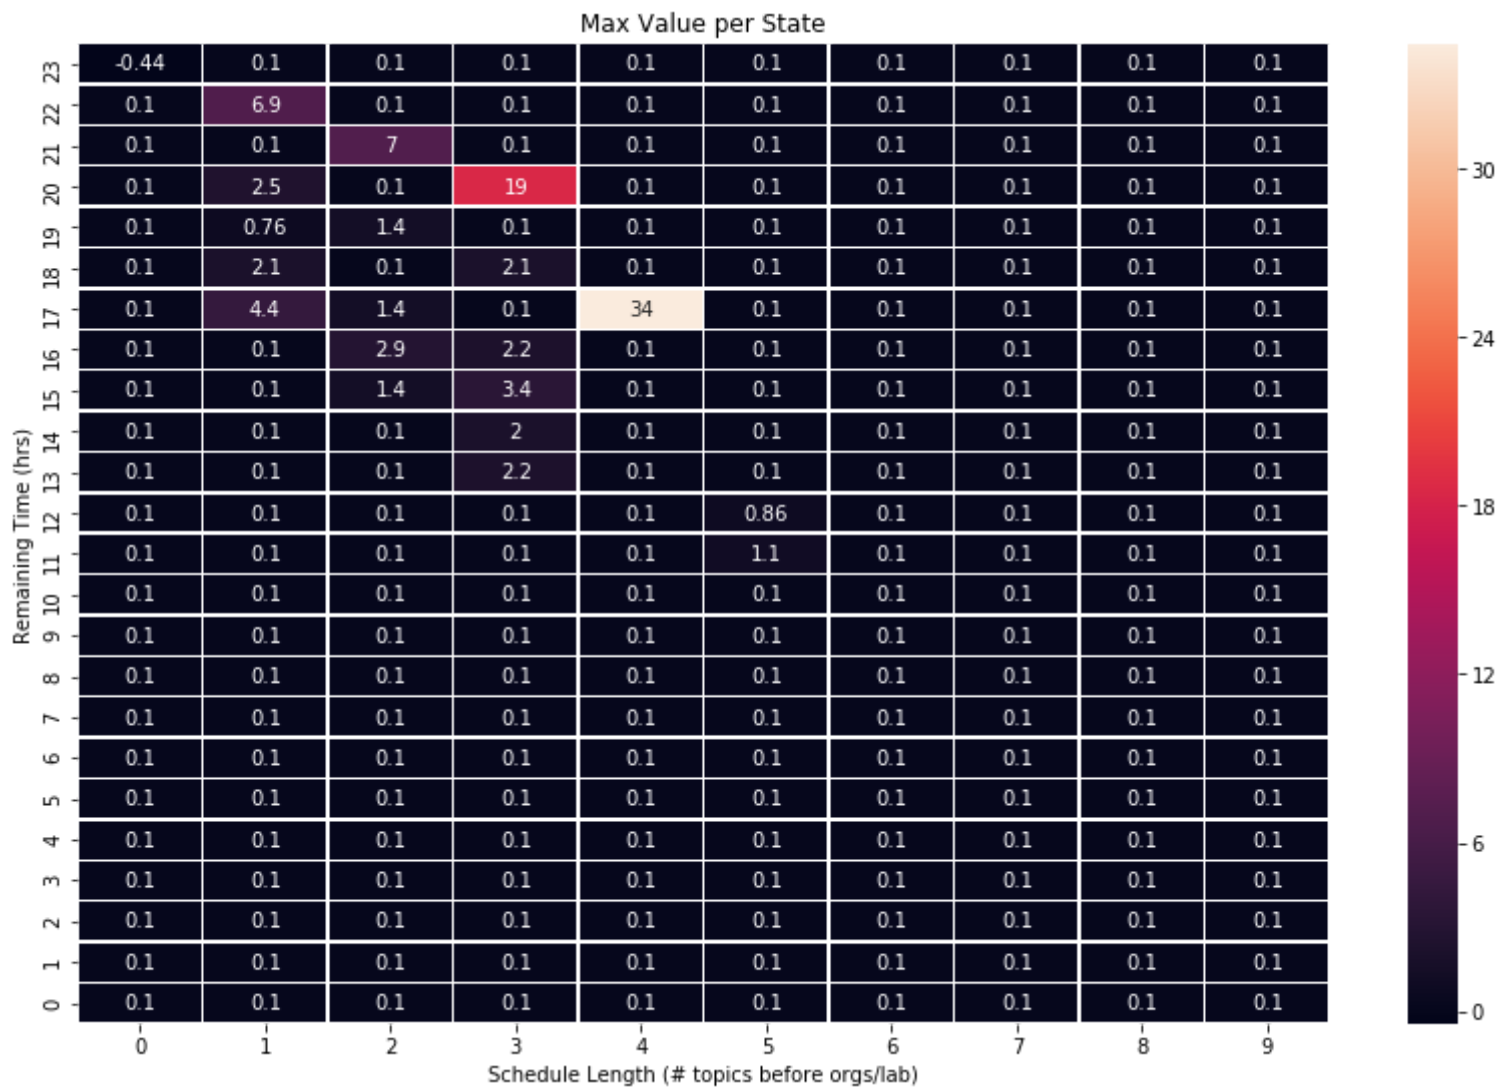
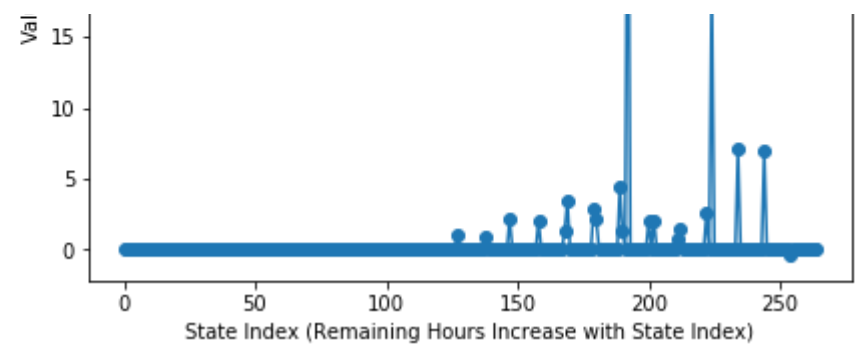
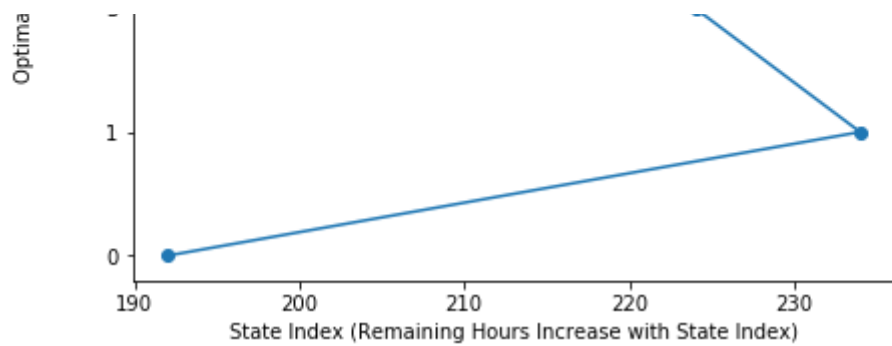
PRIOR TO ORGS/LABSTUFF:
Optimal Remaining Time: 18 HOURS
Optimal Schedule Length: 1 COMMITMENTS

WITH ORGS/LABS IN FINAL CONSIDERATION
Optimal Schedules:
Partially Optimal Schedules:
{ 'Orgs': 3, 'webtooning': 3, 'LabStuff': 1 } | Fulfillment: 0.9
{ 'Orgs': 4, 'topology': 5, 'LabStuff': 1 } | Fulfillment: 1.15
{ 'Orgs': 1, 'topology': 5, 'LabStuff': 1 } | Fulfillment: 0.875
{ 'Orgs': 4, 'animation': 3, 'LabStuff': 2 } | Fulfillment: 0.9
{ 'Orgs': 2, 'webtooning': 1, 'psych/linguistics': 5, 'LabStuff': 2 } | Fulfillment: 0.9
{ 'animation': 5, 'Orgs': 1, 'LabStuff': 2 } | Fulfillment: 0.475
{ 'Orgs': 4, 'R/statistical packages': 5, 'LabStuff': 1 } | Fulfillment: 0.875
{ 'Orgs': 2, 'LabStuff': 2, 'topology': 3 } | Fulfillment: 0.575

INIT Q-VALUES AS 0.1:

HBox(children=(IntProgress(value=0, max=100000), HTML(value='')))




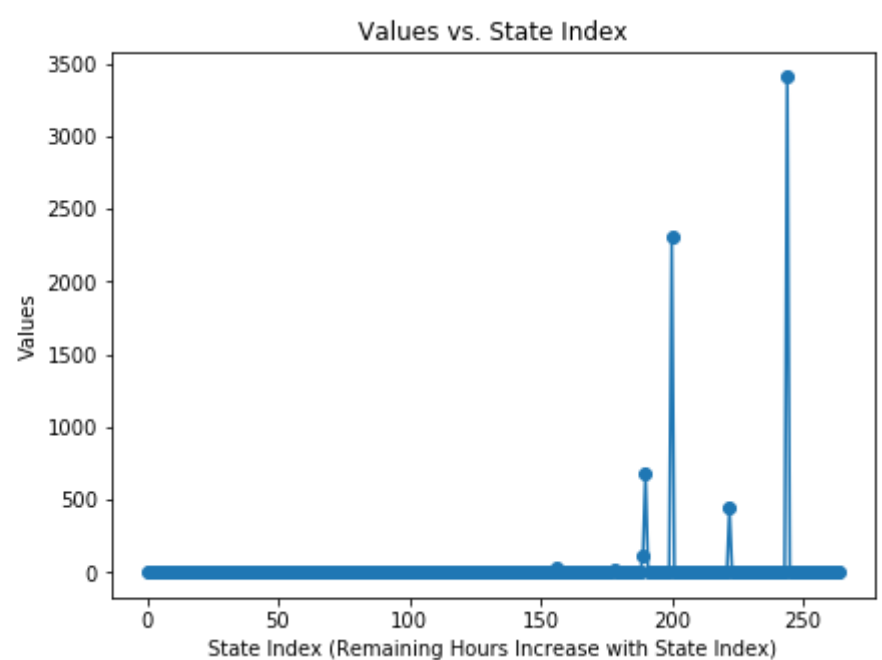



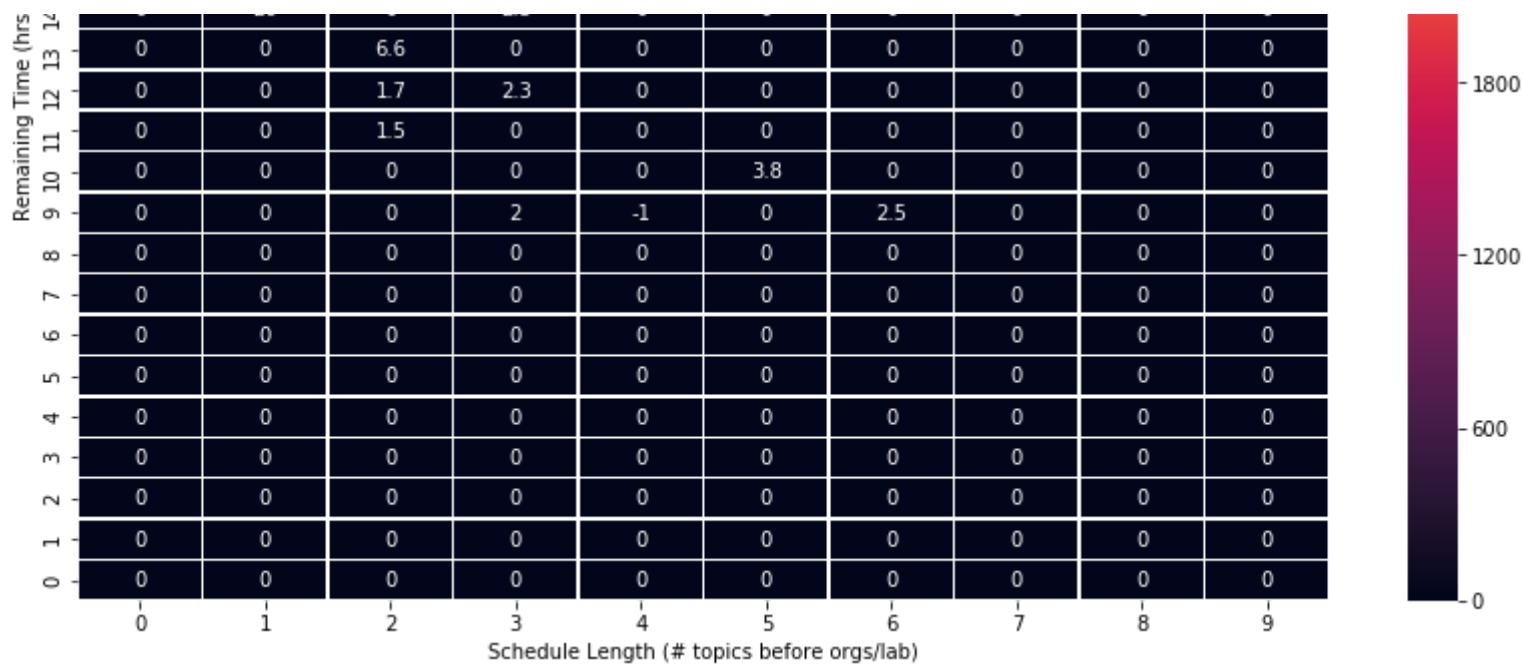
PRIOR TO ORGS/LABSTUFF:
 Optimal Remaining Time: 18 HOURS
 Optimal Schedule Length: 4 COMMITMENTS

WITH ORGS/LABS IN FINAL CONSIDERATION
 Optimal Schedules:
 {'Japanese': 3, 'LabStuff': 1, 'topology': 1, 'Orgs': 3, 'R/statistical packages': 1, 'physics': 1} | Fulfillment: 3.25

INIT Q-VALUES AS zeros:

HBox(children=(IntProgress(value=0, max=100000), HTML(value='')))



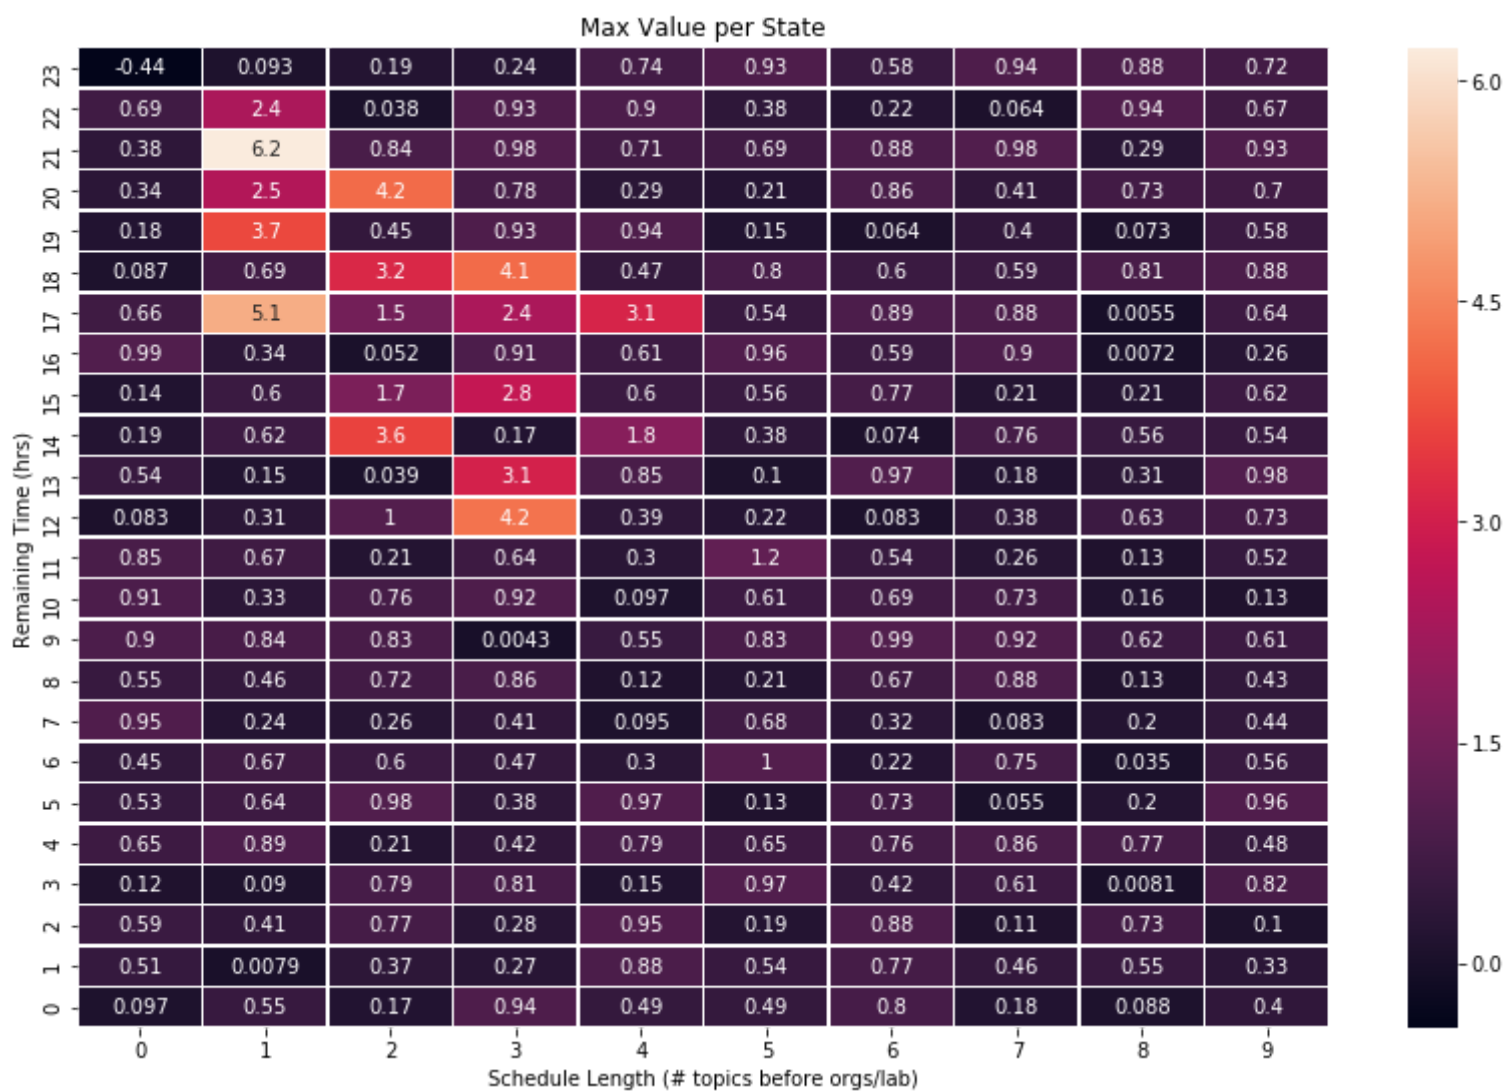
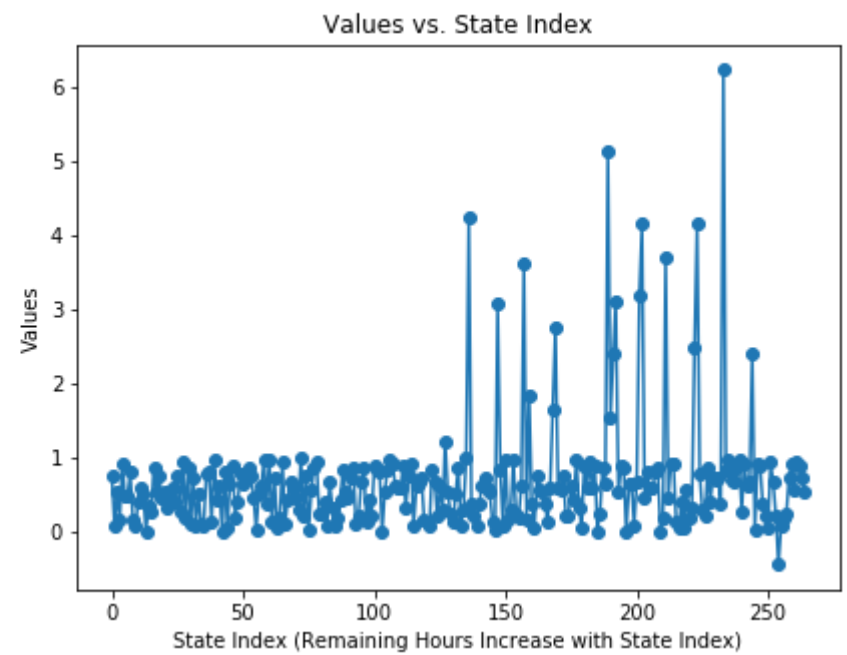
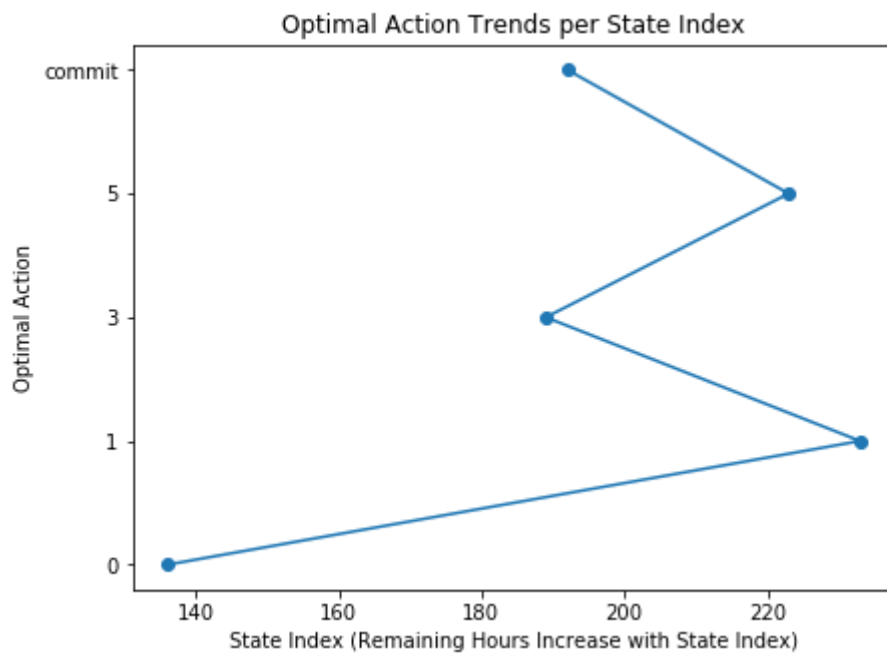


PRIOR TO ORGS/LABSTUFF:
Optimal Remaining Time: 23 HOURS
Optimal Schedule Length: 1 COMMITMENTS

WITH ORGS/LABS IN FINAL CONSIDERATION
Optimal Schedules:
{'LabStuff': 1, 'Orgs': 2, 'Japanese': 1} | Fulfillment: -0.5
{'topology': 1, 'Orgs': 2, 'LabStuff': 2} | Fulfillment: -0.5

INIT Q-VALUES AS arb:

HBox(children=(IntProgress(value=0, max=100000), HTML(value='')))



```
with open('MC_policies.pickle', 'wb') as mc_results1:  
    pickle.dump(policies, mc_results1)
```

```

with open('MC_init_qs.pickle', 'wb') as mc_results2:
    pickle.dump(MC_init_qs, mc_results2)
with open('MC_init_eps.pickle', 'wb') as mc_results3:
    pickle.dump(MC_init_eps, mc_results3)
with open('MC_init_ep_rewards.pickle', 'wb') as mc_results4:
    pickle.dump(MC_init_ep_rewards, mc_results4)
with open('MC_init_opt_actions.pickle', 'wb') as mc_results5:
    pickle.dump(MC_init_opt_actions, mc_results5)

```

▼ Algorithm 2: TD(λ)

```

# epsilon-greedy policy
def epsilon_greedy(epsilon, values, state):
    state_indx = pos_states.index(state)
    be_greedy = (np.random.random() > epsilon)
    if be_greedy:
        action = np.argmax(values[state_indx]) #optimal action
    else:
        action = np.random.choice(possible_actions)
    return action

```

```

# # initialize V(s)
def init_values(init):
    if init == "zeros":
        values = np.zeros((len(pos_states)))
    elif init == 'arb':
        values = np.ones((len(pos_states)))
        for state_indx, state in enumerate(pos_states):
            if state in terminal_states:
                values[state_indx] = 0
            else:
                values[state_indx] = np.random.random()
    else:
        values = np.ones((len(pos_states))) * init
    return values

```

```

def TD(init, lmbda, num_eps):
    # initialize V(s) arbitrarily but set to 0 if state is terminal
    values = init_values(init)

    all_episodes = []
    all_episode_rewards = []

    for episode in tqdm(range(num_eps)):

        #initialize weights
        e_weights = np.zeros(len(pos_states))

        #initialize S
        remaining_time = 24
        schedule = {}
        topic = np.random.choice(possible_topics)

        state = [remaining_time, len(schedule)]
        state_indx = pos_states.index(state)

        episode = []
        episode_rewards = []

        # for each step in episode
        while True:
            #take action, observe reward, new_state
            action = epsilon_greedy(epsilon, values, state)
            action_indx = possible_actions.index(action)

            if action == 7:
                new_state, reward, schedule, remaining_time = step(remaining_time, schedule, topic, action)
                new_state_indx = pos_states.index(new_state)
            elif action != 7:
                new_state, reward, schedule, topic = step(remaining_time, schedule, topic, action)
                new_state_indx = pos_states.index(new_state)

            ep = [state, action, reward, schedule]
            episode.append(ep)
            episode_rewards.append(reward)

            # update error and weights
            td_error = reward + gamma*values[new_state_indx] - values[state_indx]
            e_weights[state_indx] = (1 - alpha) * e_weights[state_indx] + 1 #dutch traces

            #update values and eligibility weights for all states
            values = values + alpha * td_error * e_weights
            e_weights = gamma * lmbda * e_weights

```

```

        if new_state == 0:
            break

        remaining_time, len_schedule = new_state
        state = [remaining_time, len_schedule]
        state_indx = pos_states.index(state)

    all_episodes.append(episode)
    all_episode_rewards.append(episode_rewards)

    return values, all_episodes, all_episode_rewards

```

```

%%time
gamma = 0.5 #1.0
lmbdas = [0,0.5,0.75,1]
epsilon = 0.1
alpha = 0.1
init = 'arb'
inits = [0.1, 'arb', 'zeros']

num_eps = 100000 # number of runs/eps

TD_init_values = []
TD_init_eps = []
TD_init_ep_rewards = []
TD_init_opt_actions = []

for init in inits:
    print("\033[1;38m  INIT VALUES as {} \033[0m".format(init))
    for lambda in lmbdas:
        print("\033[1;43m  LAMBDA = {} \033[0m".format(lambda))

        values, all_episodes, all_episode_rewards = TD(init, lambda, num_eps)
        optimal_actions = get_opt_actions("TD", values)

        TD_init_values.append(values)
        TD_init_eps.append(all_episodes)
        TD_init_ep_rewards.append(all_episode_rewards)
        TD_init_opt_actions.append(optimal_actions)

#plot stuff
plot_heatmap(values)
plot_values_actions_states("TD", values)
if np.argmax(values) != 0:
    best_r_time, best_len = pos_states[np.argmax(values)]
    print("PRIOR TO ORGS/LAB:")
    print("Optimal Free Time: {} HOURS \nOptimal Schedule Length (without Orgs/Lab): {} COMMITMENTS".format(best_r_time, best_len))

# returns schedules that are EITHER the optimal length or optimal, get unique entries
optimal_schedules, part_optimal_schds, opt_schd_rewards, part_opt_rewards = get_opt_schedules(values, all_episodes, optimal_actions)
optimal_schedules = [dict(entry) for entry in set(frozenset(schd.items()) for schd in optimal_schedules)]
part_optimal_schds = [dict(entry) for entry in set(frozenset(schd.items()) for schd in part_optimal_schds)]

print("FACTORING IN ORGS/LAB:")
print("Optimal Schedules:")
for indx, schd in enumerate(optimal_schedules):
    print(schd, "| Fulfillment:", opt_schd_rewards[indx])

print()
if optimal_schedules == []:
    print("Partially Optimal Schedules:")
    for indx, schd in enumerate(part_optimal_schds):
        print(schd, "| Fulfillment:", part_opt_rewards[indx])
print("-----\n")
sys.stdout.flush()

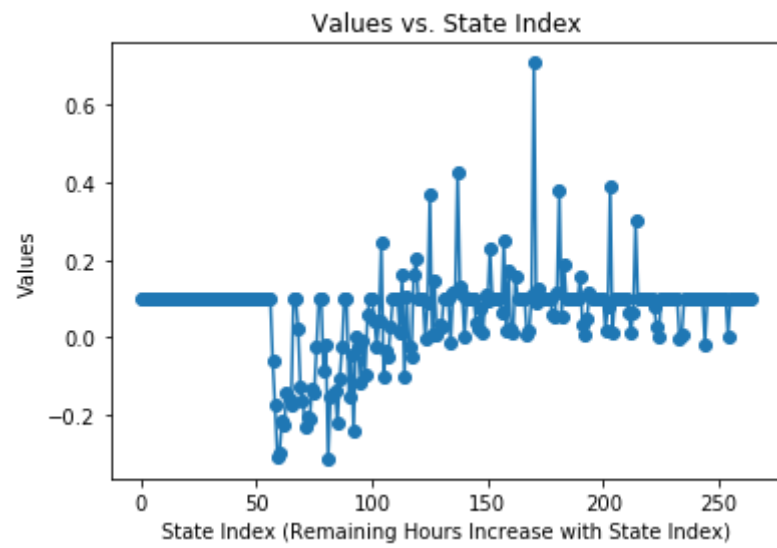
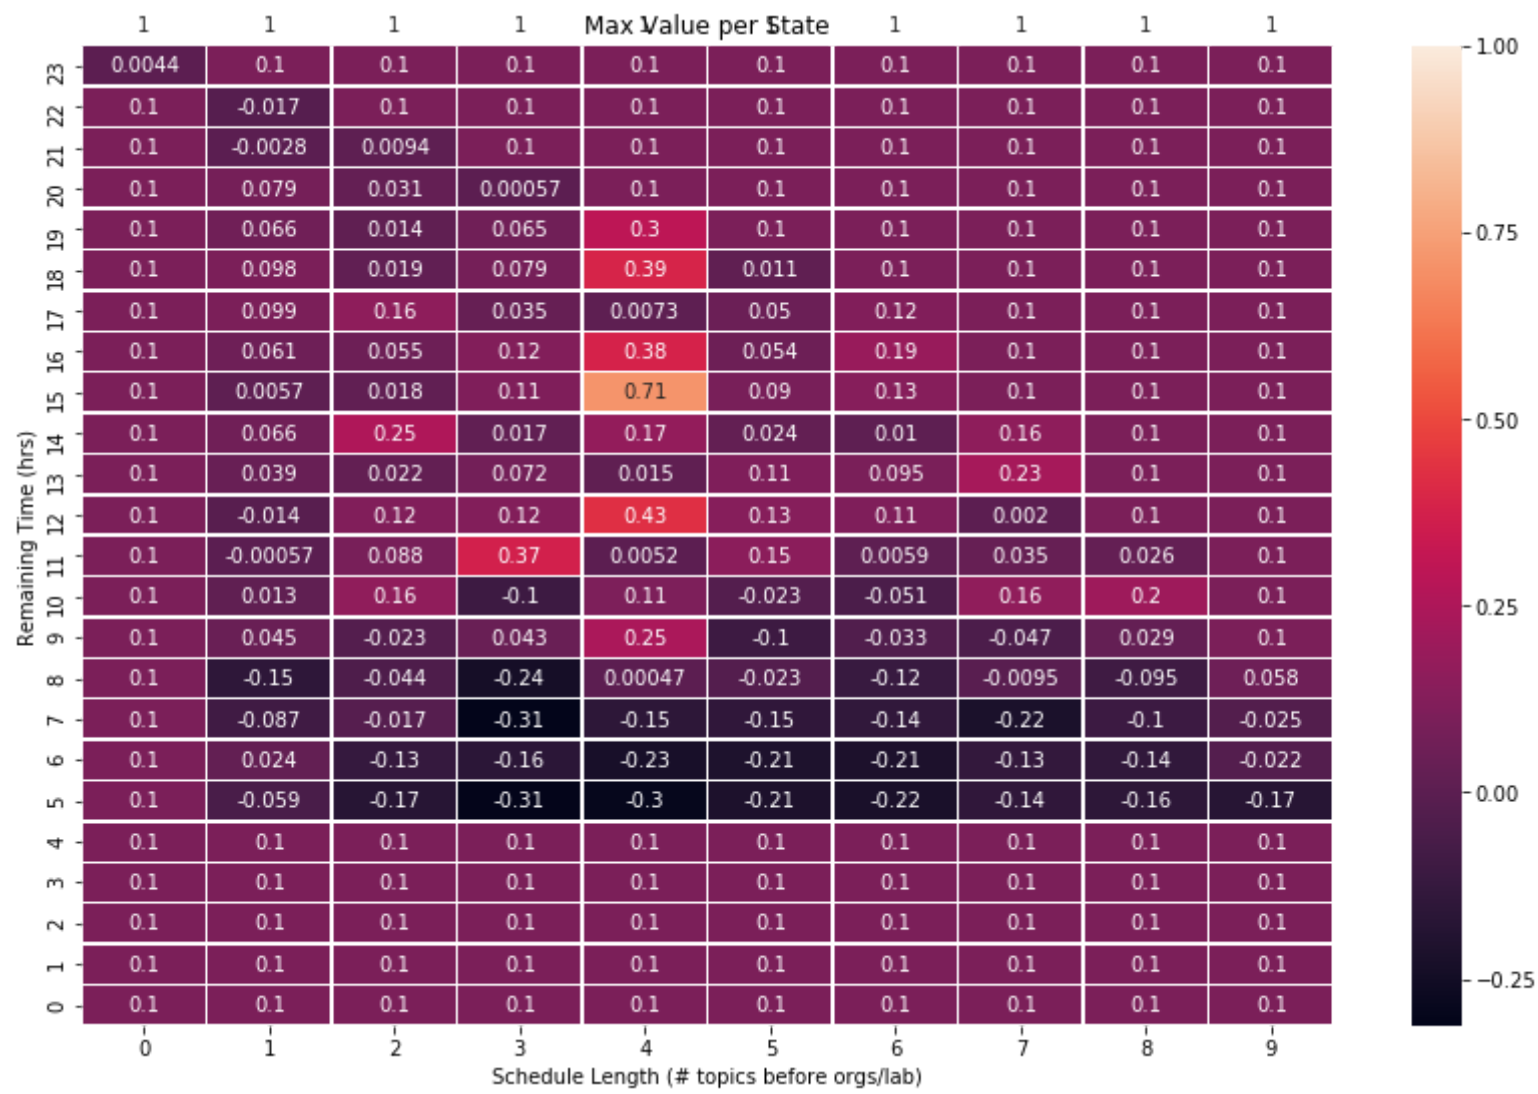
```



INIT VALUES as 0.1

LAMBDA = 0

HBox(children=(IntProgress(value=0, max=100000), HTML(value='')))



PRIOR TO ORGS/LAB:

Optimal Free Time: 16 HOURS

Optimal Schedule Length (without Orgs/Lab): 4 COMMITMENTS

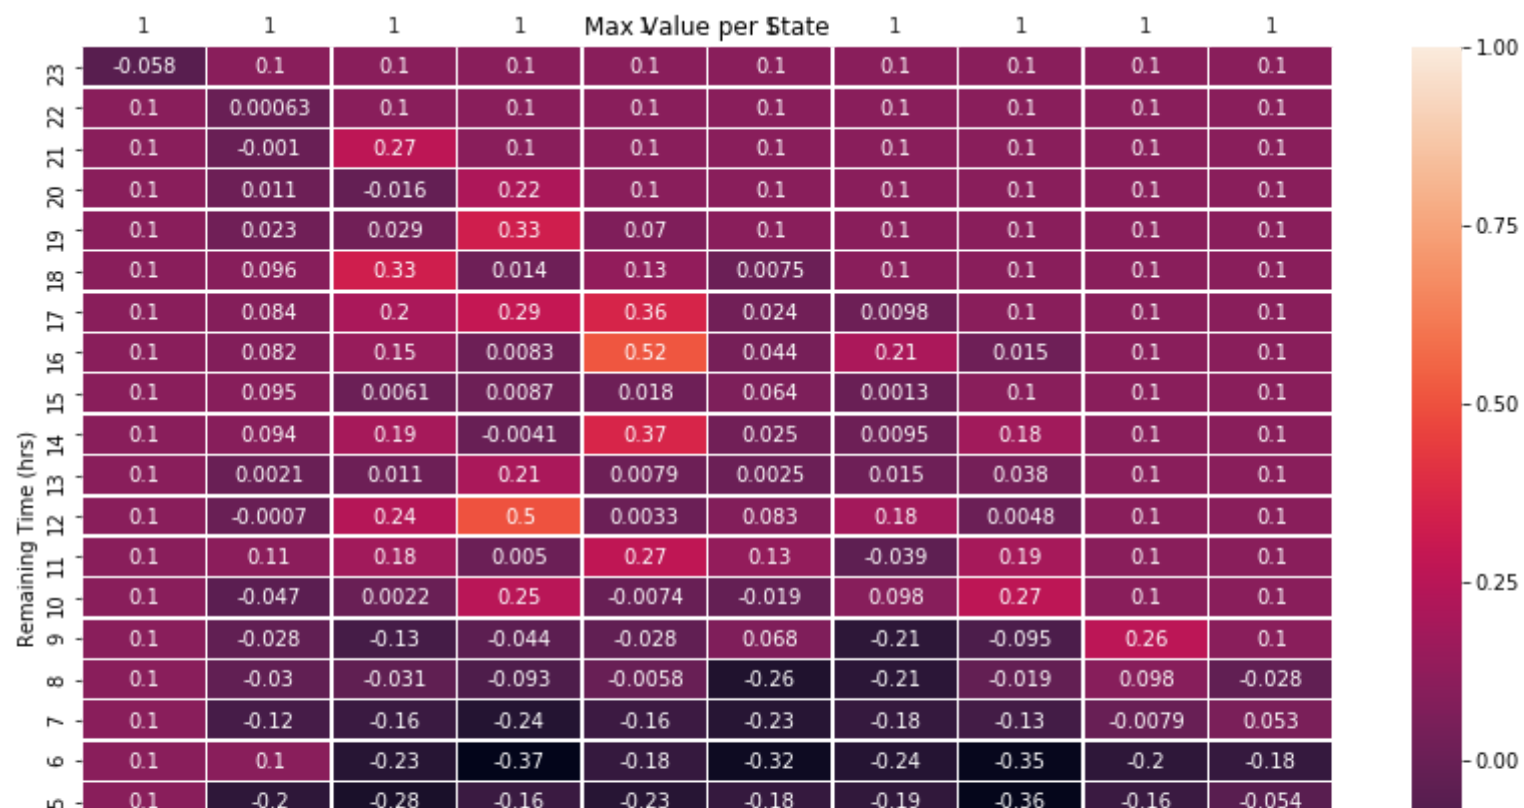
FACTORING IN ORGS/LAB:

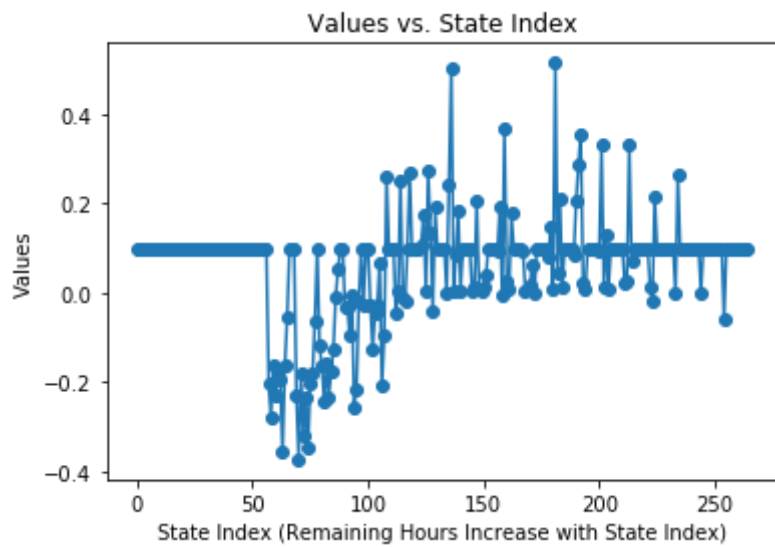
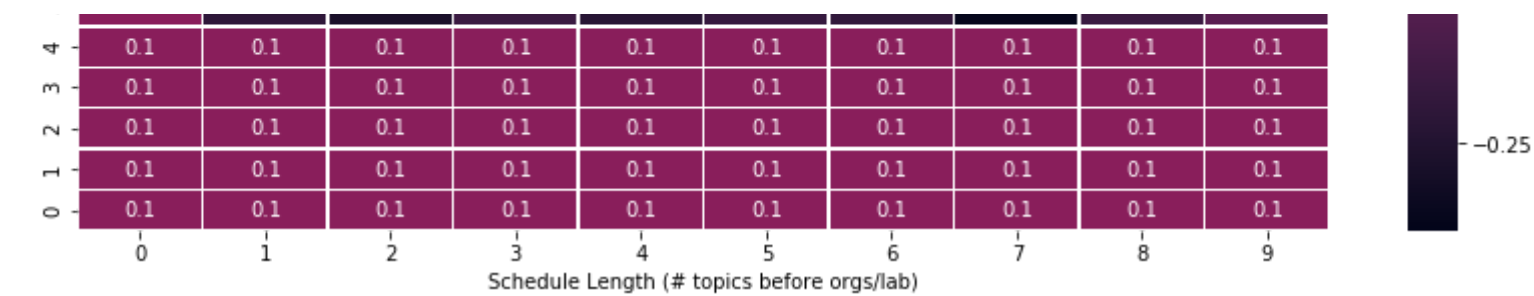
Optimal Schedules:

Partially Optimal Schedules:

LAMBDA = 0.5

HBox(children=(IntProgress(value=0, max=100000), HTML(value='')))





PRIOR TO ORGS/LAB:

Optimal Free Time: 17 HOURS

Optimal Schedule Length (without Orgs/Lab): 4 COMMITMENTS

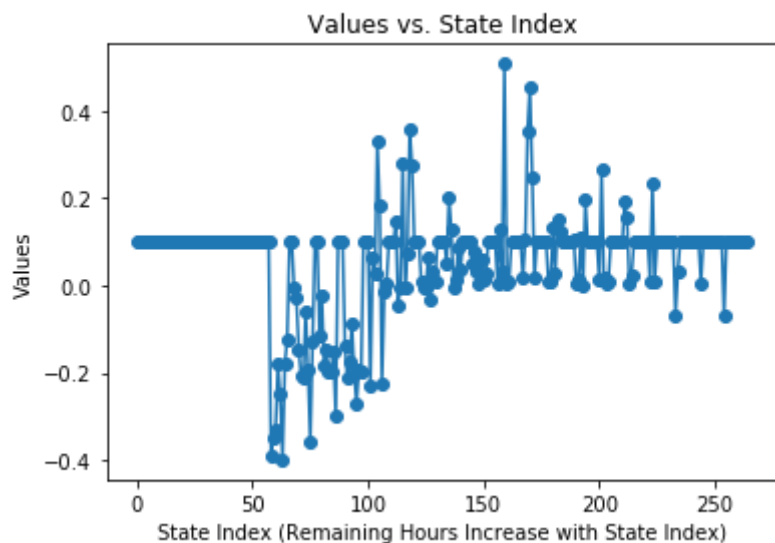
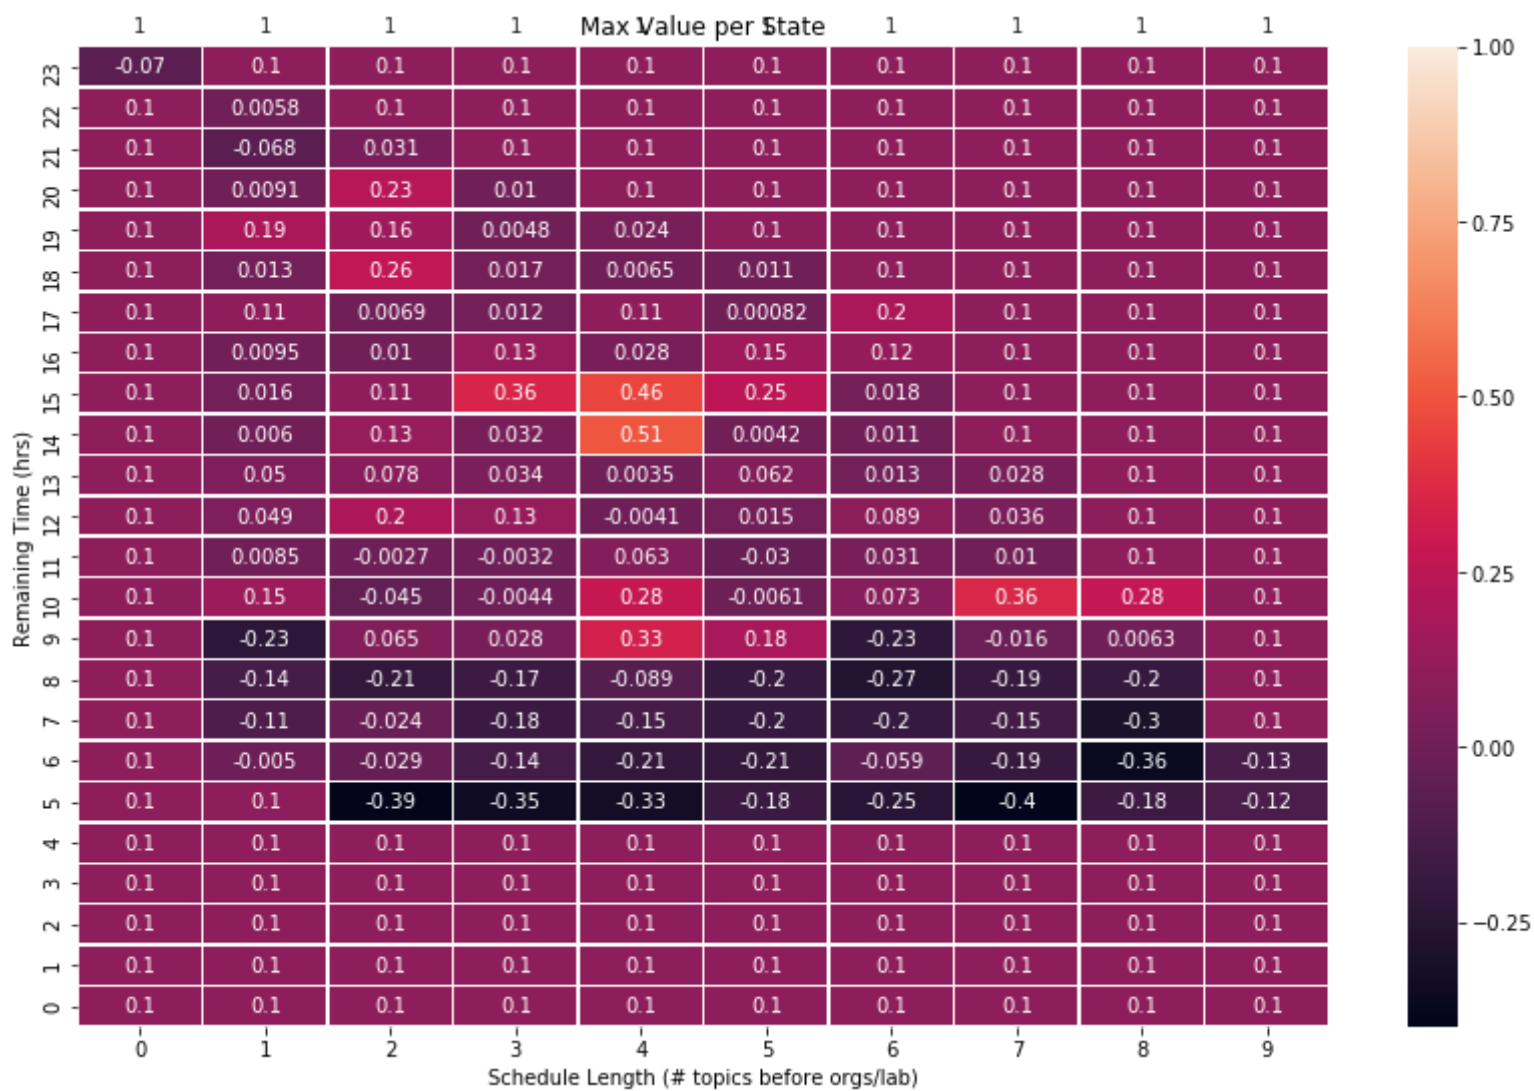
FACTORING IN ORGS/LAB:

Optimal Schedules:

Partially Optimal Schedules:

LAMBDA = 0.75

HBox(children=(IntProgress(value=0, max=100000), HTML(value='')))



PRIOR TO ORGS/LAB:

Optimal Free Time: 15 HOURS

Optimal Schedule Length (without Orgs/Lab): 4 COMMITMENTS



FACTORING IN ORGS/LAB:

Optimal Schedules:

FACTORING IN ORGS/LAB:
Optimal Schedules:

Partially Optimal Schedules:


LAMBDA = 1
HBox(children=(IntProgress(value=0, max=100000), HTML(value='')))





PRIOR TO ORGS/LAB:
Optimal Free Time: 13 HOURS
Optimal Schedule Length (without Orgs/Lab): 4 COMMITMENTS
FACTORING IN ORGS/LAB:
Optimal Schedules:

Partially Optimal Schedules:

INIT VALUES as arb
LAMBDA = 0
HBox(children=(IntProgress(value=0, max=100000), HTML(value='')))







PRIOR TO ORGS/LAB:
Optimal Free Time: 4 HOURS
Optimal Schedule Length (without Orgs/Lab): 3 COMMITMENTS
FACTORING IN ORGS/LAB:
Optimal Schedules:

Partially Optimal Schedules:

LAMBDA = 0.5
HBox(children=(IntProgress(value=0, max=100000), HTML(value='')))



0 50 100 150 200 250
State Index (Remaining Hours Increase with State Index)

PRIOR TO ORGS/LAB:

Optimal Free Time: 5 HOURS

Optimal Schedule Length (without Orgs/Lab): 2 COMMITMENTS

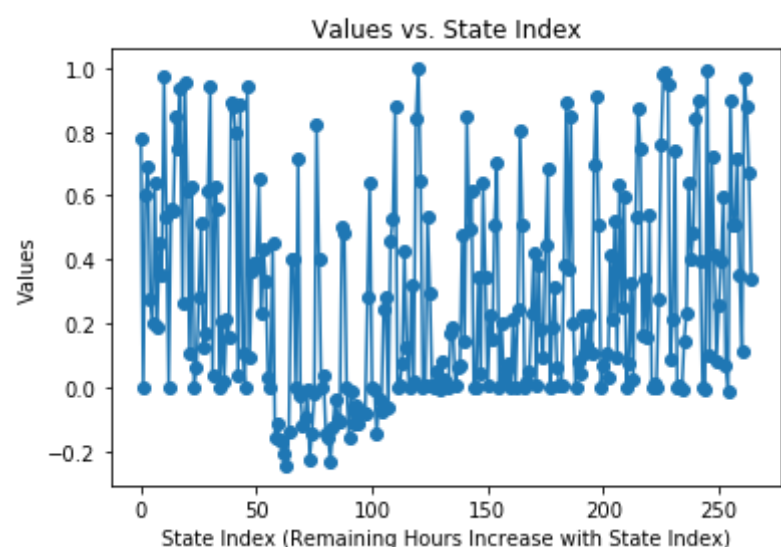

FACTORING IN ORGS/LAB:

Optimal Schedules:

Partially Optimal Schedules:

LAMBDA = 0.75

HBox(children=(IntProgress(value=0, max=100000), HTML(value='')))



PRIOR TO ORGS/LAB:

Optimal Free Time: 11 HOURS

Optimal Schedule Length (without Orgs/Lab): 9 COMMITMENTS


FACTORING IN ORGS/LAB:

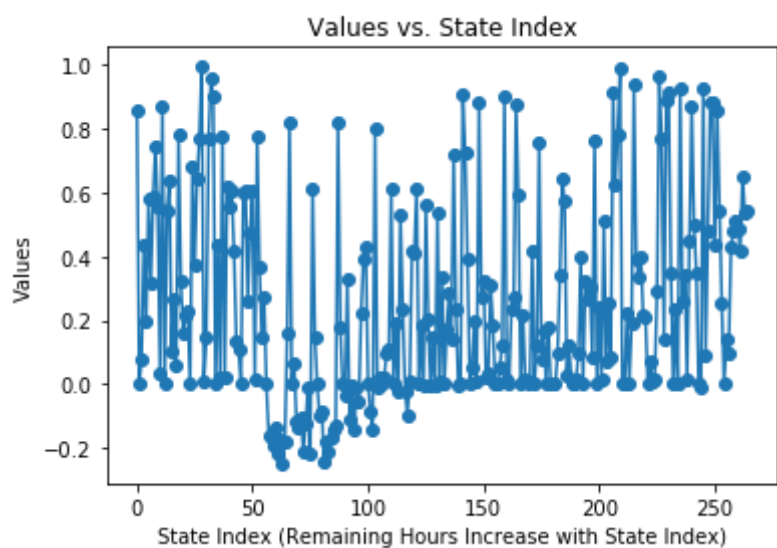
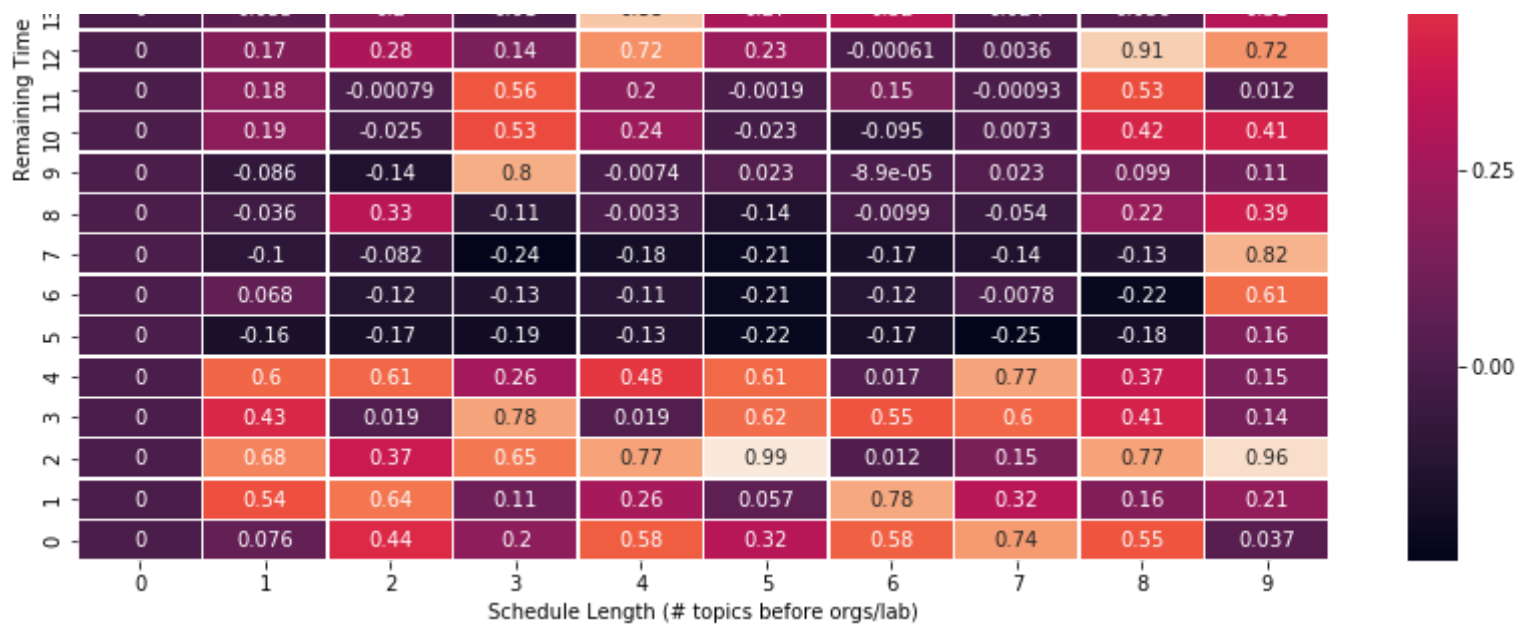
Optimal Schedules:

Partially Optimal Schedules:

LAMBDA = 1

HBox(children=(IntProgress(value=0, max=100000), HTML(value='')))





PRIOR TO ORGS/LAB:

Optimal Free Time: 3 HOURS

Optimal Schedule Length (without Orgs/Lab): 5 COMMITMENTS

FACTORING IN ORGS/LAB:



Optimal Schedules:

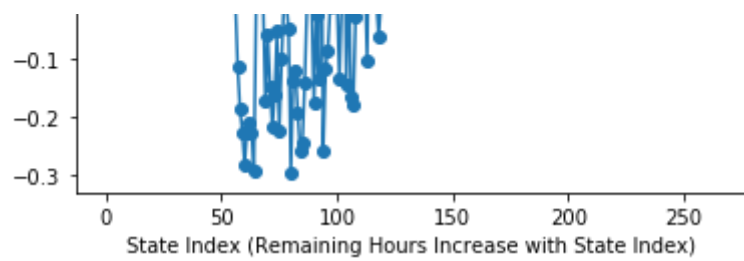
Partially Optimal Schedules:

INIT VALUES as zeros

LAMBDA = 0

HBox(children=(IntProgress(value=0, max=100000), HTML(value='')))







PRIOR TO ORGS/LAB:
 Optimal Free Time: 21 HOURS
 Optimal Schedule Length (without Orgs/Lab): 3 COMMITMENTS
 FACTORING IN ORGS/LAB:
 Optimal Schedules:

Partially Optimal Schedules:

LAMBDA = 0.5

HBox(children=(IntProgress(value=0, max=100000), HTML(value='')))



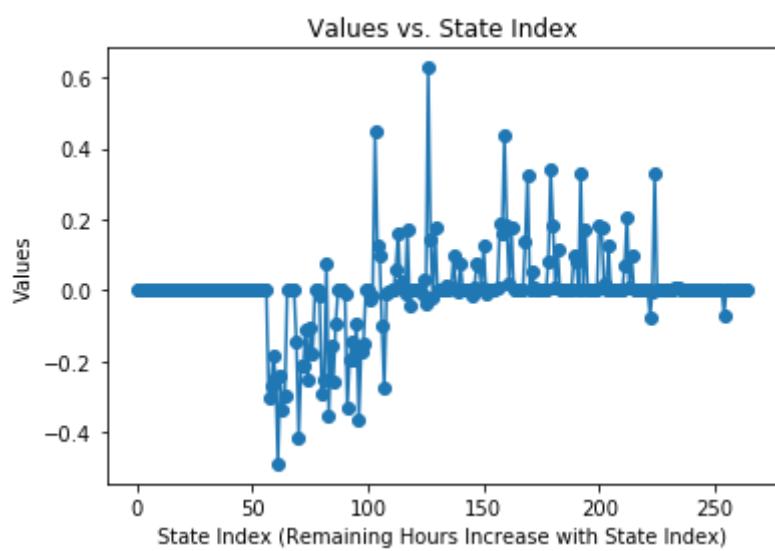
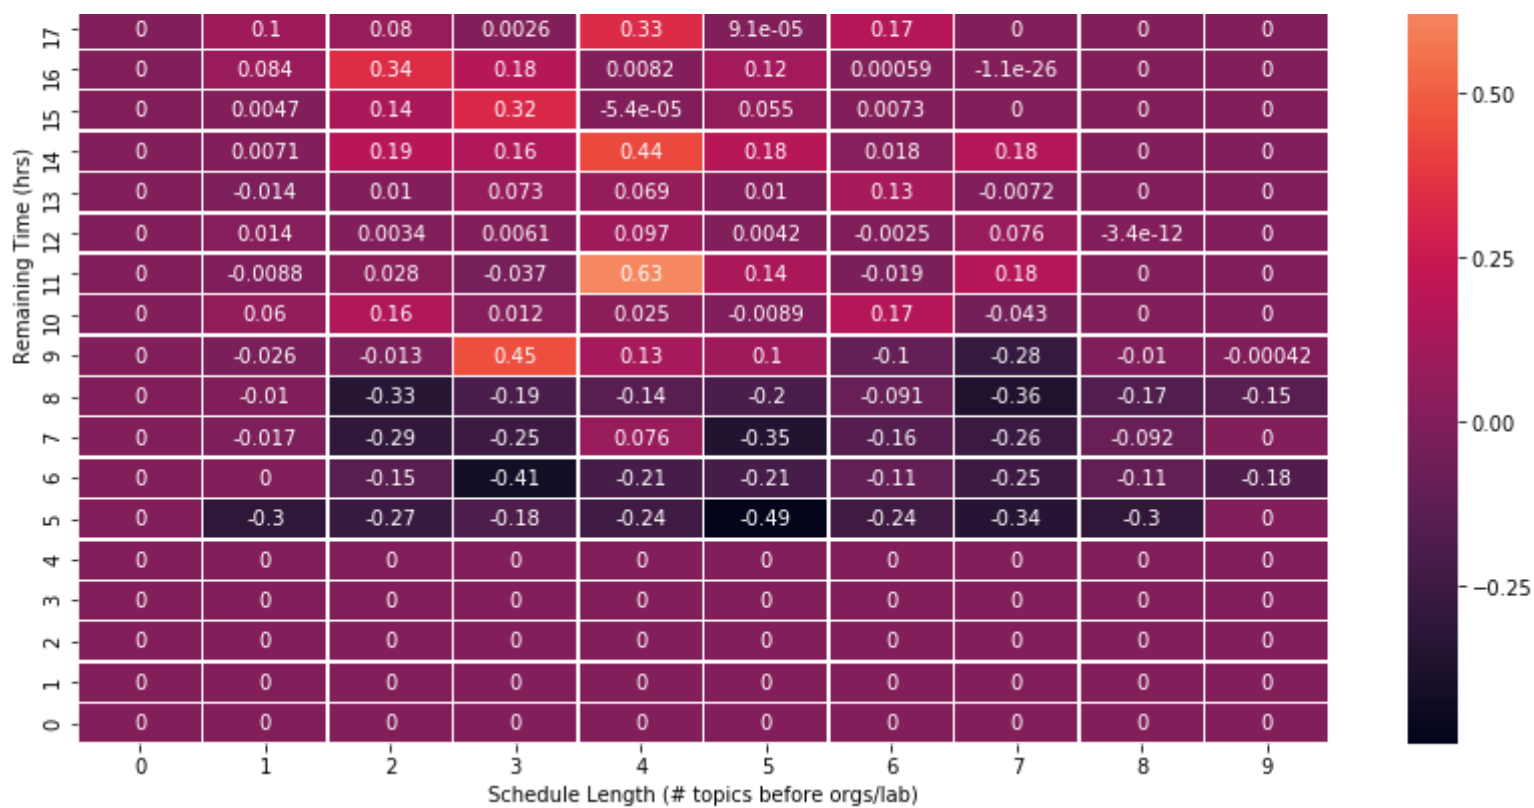
PRIOR TO ORGS/LAB:
 Optimal Free Time: 18 HOURS
 Optimal Schedule Length (without Orgs/Lab): 4 COMMITMENTS
 FACTORING IN ORGS/LAB:
 Optimal Schedules:

Partially Optimal Schedules:

LAMBDA = 0.75

HBox(children=(IntProgress(value=0, max=100000), HTML(value='')))





PRIOR TO ORGS/LAB:

Optimal Free Time: 12 HOURS

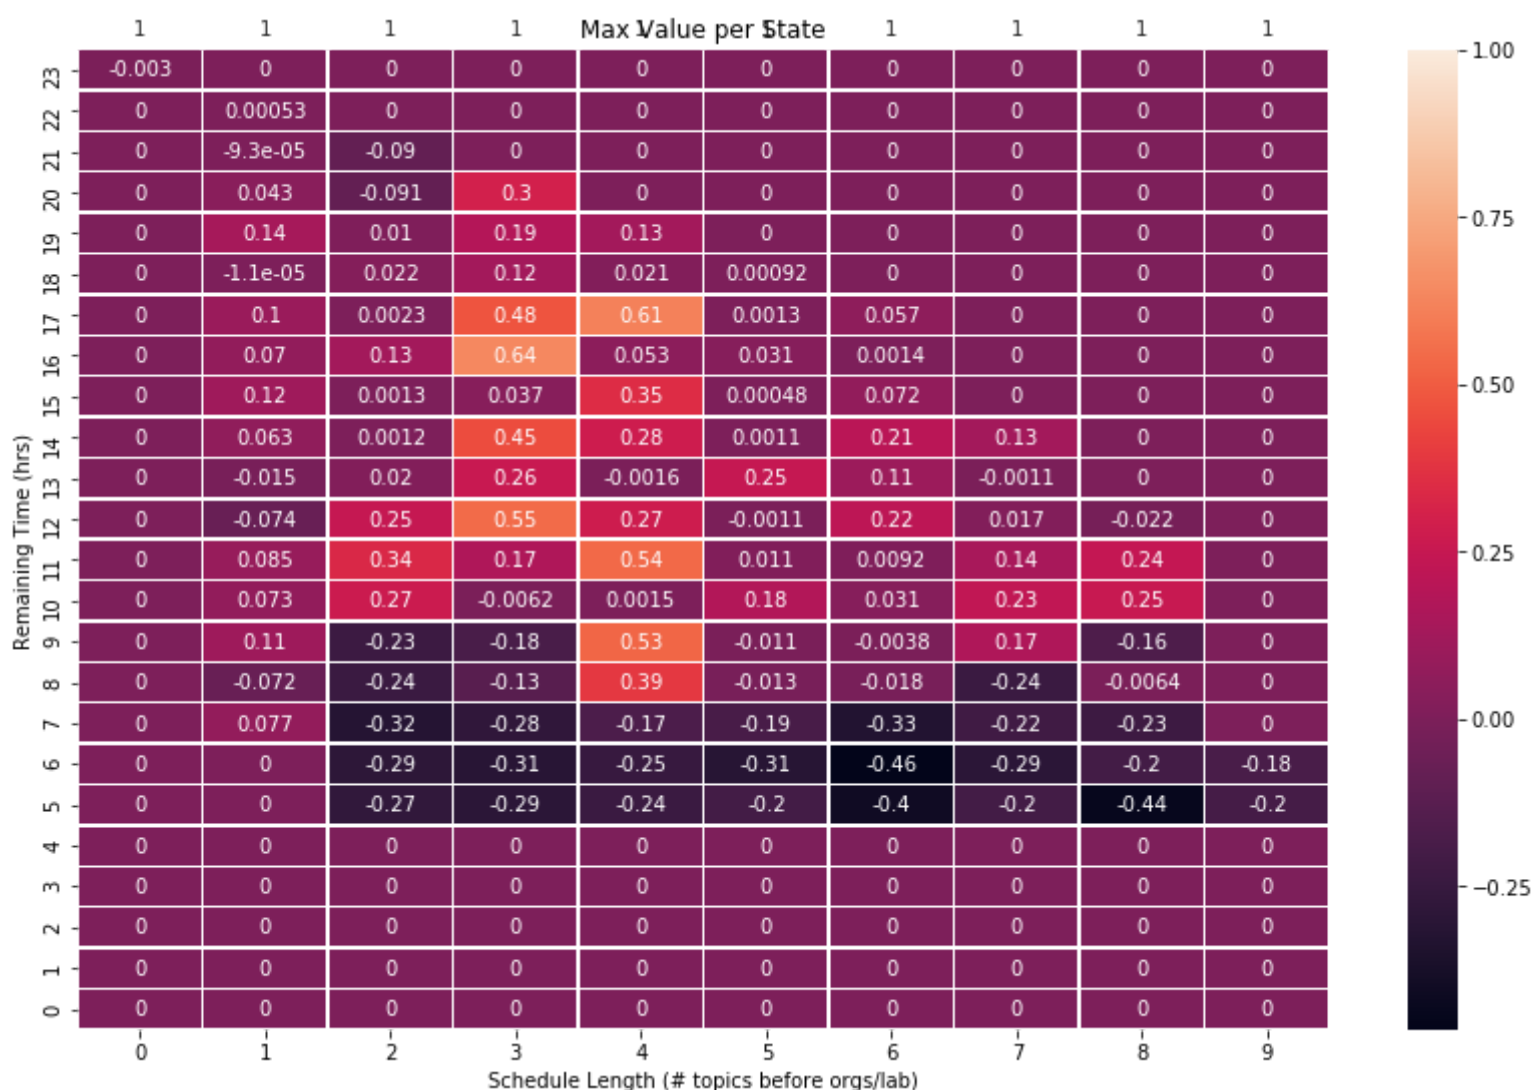
Optimal Schedule Length (without Orgs/Lab): 4 COMMITMENTS

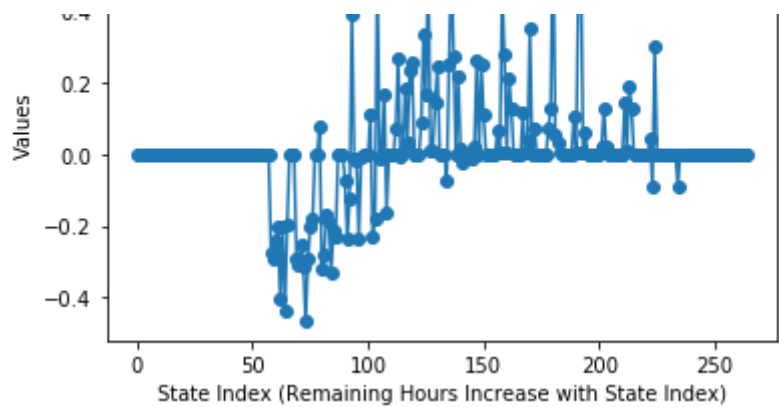
FACTORING IN ORGS/LAB:

Optimal Schedules:

Partially Optimal Schedules:

LAMBDA = 1
HBox(children=(IntProgress(value=0, max=100000), HTML(value='')))





PRIOR TO ORGS/LAB:
Optimal Free Time: 17 HOURS
Optimal Schedule Length (without Orgs/Lab): 3 COMMITMENTS
FACTORTING IN ORGS/LAB:

▼ Save Variables

```
-----  
with open('TD_init_values.pickle', 'wb') as td_results1:  
    pickle.dump(TD_init_values, td_results1)  
with open('TD_init_eps.pickle', 'wb') as td_results2:  
    pickle.dump(TD_init_eps, td_results2)  
with open('TD_init_ep_rewards.pickle', 'wb') as td_results3:  
    pickle.dump(TD_init_ep_rewards, td_results3)
```

```
with open('functions.pickle', 'wb') as functions:  
    pickle.dump([quick_check,  
                get_opt_actions,  
                calc_V_star,  
                plot_values_actions_states,  
                plot_heatmap,  
                get_opt_schedules], functions)
```